

DCS 530 : Final Project

"Are marriages made in heaven?"

Aniruddha Joshi

Date: Aug 07, 2023

Data set Reference : <https://gssdataexplorer.norc.umd.edu/> (<https://gssdataexplorer.norc.umd.edu/>)

A few months ago, one of my friends got engaged, and we celebrated with an engagement party at a nice restaurant. "It looks like your marriage is created in heaven," said one of my friends. This sentence made me wonder: Why do people say things like that? Are marriage truly created in heaven? If so, why do relationships end in divorce? Is it possible to foresee whether a couple will be happy together forever? I choose to conduct some data analysis to find out. The link to the General Social Survey information found at <https://gssdataexplorer.norc.umd.edu/adfdfs/> (<https://gssdataexplorer.norc.umd.edu/adfdfs/>).

```
In [44]: #Download the required files for execution
from os.path import basename, exists

def download(url):
    filename = basename(url)
    if not exists(filename):
        from urllib.request import urlretrieve

        local, _ = urlretrieve(url, filename)
        print("Downloaded " + local)

download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/thinkstats2.py")
download("https://github.com/AllenDowney/ThinkStats2/raw/master/code/thinkplot.py")
```

Import the required libraries

```
In [45]: #import the required libraries
import numpy as np
import pandas as pd

import matplotlib.pyplot as plt
import thinkstats2
import thinkplot
import IPython.display

import seaborn as sns

from sklearn.preprocessing import MinMaxScaler

from scipy.stats import pearsonr
from scipy.stats import chi2_contingency
import statsmodels.api as sm
from sklearn.linear_model import LogisticRegression
from sklearn.impute import SimpleImputer
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import train_test_split
```

Import the dataset

```
In [46]: #Import datasets

dataset_gss = pd.read_csv('C:\\Users\\aniruddha.joshi\\OneDrive - Emerson\\Personal\\MS Data Science Emerson\\DCS 530\\Project
```

In [47]:

A quick review of the data
dataset_gss

Out[47]:

	year	id_	hrs1	wrkslf	marital	divorce	spwrksta	childs	age	educ	spdeg	sex	race	family16	income	re
0	1972	1	.i: Inapplicable	Someone else	Never married	.i: Inapplicable	.i: Inapplicable	0	23	16	.i: Inapplicable	FEMALE	White	FATHER	.i: Inapplicable	Inappli
1	1972	2	.i: Inapplicable	Someone else	Married	NO	KEEPING HOUSE	5	70	10	HIGH SCHOOL	MALE	White	M AND F RELATIVES	.i: Inapplicable	Inappli
2	1972	3	.i: Inapplicable	Someone else	Married	NO	WORKING FULLTIME	4	48	12	.n: No answer	FEMALE	White	MOTHER & FATHER	.i: Inapplicable	Inappli
3	1972	4	.i: Inapplicable	Someone else	Married	NO	WORKING FULLTIME	0	27	17	GRADUATE	FEMALE	White	MOTHER & FATHER	.i: Inapplicable	Inappli
4	1972	5	.i: Inapplicable	Someone else	Married	NO	TEMP NOT WORKING	2	61	12	HIGH SCHOOL	FEMALE	White	MOTHER & FATHER	.i: Inapplicable	Inappli
...
72385	2022	3541	48	Someone else	Never married	.i: Inapplicable	.x: Not available in this release	0	22	12	.x: Not available in this release	FEMALE	White	MOTHER & STPFATHER	\$25,000 or more	.) availa this re
72386	2022	3542	50	Someone else	Married	YES	.x: Not available in this release	2	29	19	.x: Not available in this release	FEMALE	White	MOTHER & FATHER	\$25,000 or more	.) availa this re
72387	2022	3543	38	Someone else	Never married	.i: Inapplicable	.x: Not available in this release	1	32	15	.x: Not available in this release	MALE	White	MOTHER & STPFATHER	\$25,000 or more	.) availa this re
72388	2022	3544	40	Someone else	Married	NO	.x: Not available in this release	0	49	17	.x: Not available in this release	FEMALE	White	MOTHER & FATHER	\$25,000 or more	.) availa this re
72389	2022	3545	40	Someone else	Married	NO	.x: Not available in this release	1	50	20	.x: Not available in this release	MALE	White	MOTHER & FATHER	\$25,000 or more	.) availa this re

72390 rows × 19 columns

In [48]:

Listing the available columns
data=dataset_gss.columns
df = pd.DataFrame(data)
df

Out[48]:

	0
0	year
1	id_
2	hrs1
3	wrkslf
4	marital
5	divorce
6	spwrksta
7	childs
8	age
9	educ
10	spdeg
11	sex
12	race
13	family16
14	income
15	relig16
16	hapmar
17	ballot
18	incomeUSD

Vailable Definations

To find out if the couple will stay together happily forever lets analyze different sets of parametersm. Below is the explanation of each variable

```

-ballot : ballot used for interview
-hrs1 : number of hours worked last week. To check if couples have a time for each other.
-wrkslf : self-emp or works for somebody. To check if they work for someone else.
-marital : marital status. The dependent variable 1.
-divorce : ever been divorced or separated. The dependent variable 2.
-spwrksta: spouse labor force status
-childs : number of children. Checking the number of childrens helps couple to stay together.
-age : age of respondent. Is there any age factor in divorce.
-educ : highest year of school completed. Is education playing any role for couple to stay together.
-spdeg : spouse's highest degree.
-sex : respondents sex. Are more males divorced than females?
-race : race of respondent. Are divorce % larger in any specific race?
-family16: living with parents when 16 yrs old. Staying long with parents helps achieve family bonding?
-income : total family income. Is there a role of income in happy marriages?
-relig16 : religion in which raised. Do certain religions having lesser % of divorce?
-hapmar : happiness of marriage. Dependant variable 3.
-year : GSS year for this respondent
-incomUSD: Income in USD
-id : Respondent id number

```

Clean up the Database

```

In [49]: #Clean up the database
#Replaced invalid/incorrect strings with NaN
#This makes replacement in the complete database
dataset_gss.replace('.i: Inapplicable', np.nan, inplace=True)
dataset_gss.replace('.d: Do not Know/Cannot Choose', np.nan, inplace=True)
dataset_gss.replace('.f: Missing Birthdate Information', np.nan, inplace=True)
dataset_gss.replace('.i: Inapplicable', np.nan, inplace=True)
dataset_gss.replace('.j: I do not have a job', np.nan, inplace=True)
dataset_gss.replace('.m: DK, NA, IAP', np.nan, inplace=True)
dataset_gss.replace('.n: No answer', np.nan, inplace=True)
dataset_gss.replace('.p: Not applicable (I have not faced this decision)/Not imputable', np.nan, inplace=True)
dataset_gss.replace('.q: Not imputable', np.nan, inplace=True)
dataset_gss.replace('.r: Refused', np.nan, inplace=True)
dataset_gss.replace('.s: Skipped on Web', np.nan, inplace=True)
dataset_gss.replace('.u: Uncodable', np.nan, inplace=True)
dataset_gss.replace('.x: Not available in this release', np.nan, inplace=True)
dataset_gss.replace('.y: Not available in this year', np.nan, inplace=True)
dataset_gss.replace('.z: Variable-specific reserve code', np.nan, inplace=True)
dataset_gss.replace('8 or more', int(8), inplace=True)
dataset_gss.replace('89 or older', int(89), inplace=True)
dataset_gss.replace('No formal schooling', int(0), inplace=True)

```

Converting the columns to appropriate data format

```

In [50]: #Convert fields to appropriate format
dataset_gss['hrs1'] = pd.to_numeric(dataset_gss['hrs1'], errors='coerce')
dataset_gss['wrkslf'] = dataset_gss['wrkslf'].astype('category')
dataset_gss['marital'] = dataset_gss['marital'].astype('category')
dataset_gss['divorce'] = dataset_gss['divorce'].astype('category')
dataset_gss['spwrksta'] = dataset_gss['spwrksta'].astype('category')
dataset_gss['childs'] = pd.to_numeric(dataset_gss['childs'], errors='coerce')
dataset_gss['age'] = pd.to_numeric(dataset_gss['age'], errors='coerce')
dataset_gss['educ'] = pd.to_numeric(dataset_gss['educ'], errors='coerce')
dataset_gss['spdeg'] = dataset_gss['spdeg'].astype('category')
dataset_gss['sex'] = dataset_gss['sex'].astype('category')
dataset_gss['race'] = dataset_gss['race'].astype('category')
dataset_gss['family16'] = dataset_gss['family16'].astype('category')
dataset_gss['income'] = dataset_gss['income'].astype('category')
dataset_gss['relig16'] = dataset_gss['relig16'].astype('category')
dataset_gss['hapmar'] = dataset_gss['hapmar'].astype('category')
dataset_gss['incomeUSD'] = pd.to_numeric(dataset_gss['incomeUSD'], errors='coerce')
dataset_gss['ballot'] = dataset_gss['ballot'].astype('string')

```

In [51]:

```
#Quick review of numeric fields
#descriptive characteristics about the variables: Mean, Mode, Spread, and Tail
dataset_gss.describe()
```

Out[51]:

	year	id_	hrs1	childs	age	educ	incomeUSD
count	72390.000000	72390.000000	41266.000000	72129.000000	71621.000000	72127.000000	63439.000000
mean	1997.715541	1241.796395	40.843285	1.916538	46.555982	13.034633	25948.477378
std	15.109995	912.273245	13.584545	1.759511	17.600417	3.182372	9403.188921
min	1972.000000	1.000000	0.000000	0.000000	18.000000	0.000000	1011.392793
25%	1985.000000	534.000000	37.000000	0.000000	32.000000	12.000000	19058.115660
50%	1998.000000	1083.000000	40.000000	2.000000	44.000000	12.000000	30941.325100
75%	2010.000000	1722.000000	48.000000	3.000000	60.000000	16.000000	33203.877585
max	2022.000000	4510.000000	88.000000	8.000000	89.000000	20.000000	35499.895520

In [94]:

```
from tabulate import tabulate
data_calc = dataset_gss[['hrs1', 'age', 'educ', 'incomeUSD', 'childs']]
mean = data_calc.mean()
mode = data_calc.mode().iloc[0] # Get the first row as mode
spread = data_calc.describe().loc[['mean', 'std', 'min', '25%', '50%', '75%', 'max']]
tails = data_calc.quantile([0.05, 0.95])

# Create a DataFrame to display the results in a tabular format
result_table = pd.DataFrame({
    'Statistic': ['Mean', 'Mode', 'Standard Deviation', 'Minimum', '25th Percentile', 'Median (50th Percentile)', '75th Perce
    'Hours Worked (hrs1)': [mean['hrs1'], mode['hrs1'], spread.loc['std', 'hrs1'], spread.loc['min', 'hrs1'], spread.loc['25%
    'Age': [mean['age'], mode['age'], spread.loc['std', 'age'], spread.loc['min', 'age'], spread.loc['25%', 'age'], spread.loc
    'Education (educ)': [mean['educ'], mode['educ'], spread.loc['std', 'educ'], spread.loc['min', 'educ'], spread.loc['25%',
    'Income (USD)': [mean['incomeUSD'], mode['incomeUSD'], spread.loc['std', 'incomeUSD'], spread.loc['min', 'incomeUSD'], spr
    'Number of Children (childs)': [mean['childs'], mode['childs'], spread.loc['std', 'childs'], spread.loc['min', 'childs'],
})

# Print the result table with headers
result_table
```

Out[94]:

	Statistic	Hours Worked (hrs1)	Age	Education (educ)	Income (USD)	Number of Children (childs)
0	Mean	40.843285	46.555982	13.034633	25948.477378	1.916538
1	Mode	40.000000	30.000000	12.000000	32322.542030	0.000000
2	Standard Deviation	13.584545	17.600417	3.182372	9403.188921	1.759511
3	Minimum	0.000000	18.000000	0.000000	1011.392793	0.000000
4	25th Percentile	37.000000	32.000000	12.000000	19058.115660	0.000000
5	Median (50th Percentile)	40.000000	44.000000	12.000000	30941.325100	2.000000
6	75th Percentile	48.000000	60.000000	16.000000	33203.877585	3.000000
7	Maximum	88.000000	89.000000	20.000000	35499.895520	8.000000
8	5th Percentile	15.000000	22.000000	8.000000	7085.611358	0.000000
9	95th Percentile	64.000000	78.000000	18.000000	35034.412535	5.000000

```
In [52]: # Print unique categorical values for each column
unique_values = {}
for column in dataset_gss.columns:
    if dataset_gss[column].dtype == 'category':
        unique_values[column] = dataset_gss[column].cat.categories.tolist()

max_len = max(len(val) for val in unique_values.values())
for col in unique_values:
    unique_values[col] += [' '] * (max_len - len(unique_values[col]))

strPrint=(pd.DataFrame(unique_values).transpose())

strPrint
```

Out[52]:

	0	1	2	3	4	5	6	7	8	9
wrkslf	Self-employed	Someone else								
marital	Divorced	Married	Never married	Separated		Widowed				
divorce	NO	YES								
spwrksta	KEEPING HOUSE	OTHER	RETIRED	SCHOOL	TEMP NOT WORKING	UNEMPL, LAID OFF	WORKING FULLTIME	WORKING PARTTIME		
spdeg	ASSOCIATE/JUNIOR COLLEGE	BACHELOR	GRADUATE	HIGH SCHOOL	LT HIGH SCHOOL					
sex	FEMALE	MALE								
race	Black	Other	White							
family16	FATHER	FATHER & STPMOTHER	FEMALE RELATIVE	M AND F RELATIVES	MALE RELATIVE	MOTHER	MOTHER & FATHER	MOTHER & STPFATHER	OTHER	
income	1, 000to2,999	10, 000to 14,999	15, 000to 19,999	20, 000to 24,999	\$25,000 or more	3, 000to 3,999	4, 000to4,999	5, 000to 5,999	6, 000to 6,999	7, 000to 7,999
relig16	BUDDHISM	CATHOLIC	CHRISTIAN	HINDUISM	INTER-NONDENOMINATIONAL	JEWISH	MUSLIM/ISLAM	NATIVE AMERICAN	NONE	ORTHODOX-CHRISTIAN
hapmar	NOT TOO HAPPY	PRETTY HAPPY	VERY HAPPY							

Plotting Histograms of general polulation




```
In [85]: #Index(['year', 'id_', 'hrs1', 'wrkslf', 'marital', 'divorce', 'spwrksta',_
#         'childs', 'age', 'educ', 'spdeg', 'sex', 'race', 'family16', 'income',
#         'relig16', 'hapmar', 'ballot', 'hsr1'],
#         dtype='object')




# Plot histograms of the height, weight, and age data
#dataset_gss[['hrs1', 'wrkslf', 'marital', 'divorce']].hist(figsize=(5, 5))
# Add mean and median lines




for column in ['hrs1', 'age', 'educ', 'incomeUSD', 'childs', 'marital', 'divorce', 'spwrksta',
               'sex', 'race', 'family16', 'income',
               'hapmar', 'wrkslf']:
    fig, ax = plt.subplots()
    dataset_gss[column].hist(ax=ax, bins=8, figsize=(12, 3))
    plt.title(f"Distribution of {column}")



    if is_numeric(dataset_gss[column]):
        ax.axvline(dataset_gss[column].mean(), color='r', linestyle='dashed', label='Mean', linewidth = 4)
        ax.axvline(dataset_gss[column].median(), color='g', linestyle='dashed', label='Median')

    handles, labels = ax.get_legend_handles_labels()
    fig.legend(handles, labels, loc="upper right")
    # Rotate x-axis labels by 30 degrees
    plt.xticks(rotation=30)
    plt.show()
```









Identifying the outliers

```
In [54]: #Plotting the outliers of 5 numeric variables

# Load the dataset
dataset = dataset_gss[['hrs1', 'age', 'educ', 'incomeUSD', 'childs']]

df = dataset

scaler = MinMaxScaler()
df_normalized = pd.DataFrame(scaler.fit_transform(df), columns=df.columns)

# Create box plots for each normalized numerical variable
plt.figure(figsize=(25, 5))
sns.set(style="whitegrid")
plt.subplot(1, 2, 2)
box_plot = sns.boxplot(data=df_normalized, palette="Set2")
plt.title('Outliers')

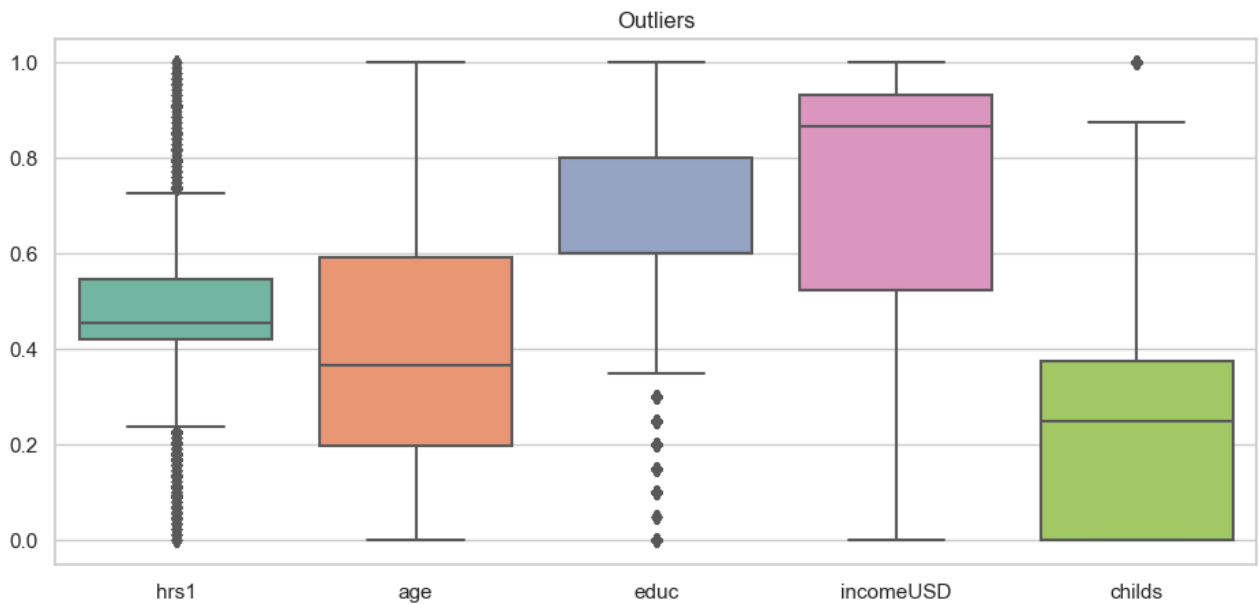
# Adding Legend
# Create a custom legend
legend_labels = ['hrs1', 'age', 'educ', 'incomeUSD', 'childs']
legend_handles = [plt.Line2D([0], [0], marker='o', color='w', markerfacecolor=sns.color_palette("Set2")[i], markersize=10) for i in range(5)]

plt.legend(legend_handles, legend_labels, loc='right')

plt.show()

# Create scatter plots for each normalized numerical variable
# plt.subplot(1, 2, 2)
# sns.pairplot(dataset_gss)
# plt.title('Scatter Plot of Normalized Numerical Variables')

# plt.tight_layout()
# plt.show()
```



Plotting histogram of divorced population

```
In [55]: # Filtering the dataset based on fivorce status
ds_divorced = dataset_gss[dataset_gss.divorce == 'YES']
```

In [56]:




```




for column in ['hrs1', 'age', 'educ', 'incomeUSD', 'childs', 'spwrksta',
               'sex', 'race', 'family16', 'wrkslf']:
    fig, ax = plt.subplots()
    ds_divorced[column].hist(ax=ax, bins=10, figsize=(12, 3))
    plt.title(f"Distribution of {column}")

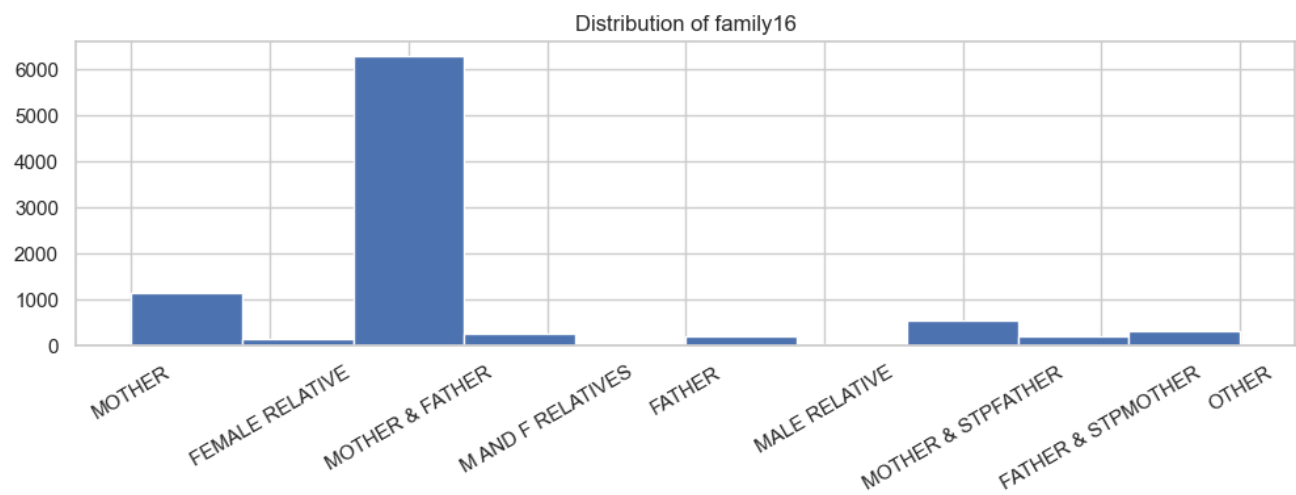

    # Rotate x-axis Labels by 30 degrees
    plt.xticks(rotation=30)

    plt.show()
    #hist = thinkstats2.Hist(ds_divorced[column])
    #thinkplot.Hist(hist)

```







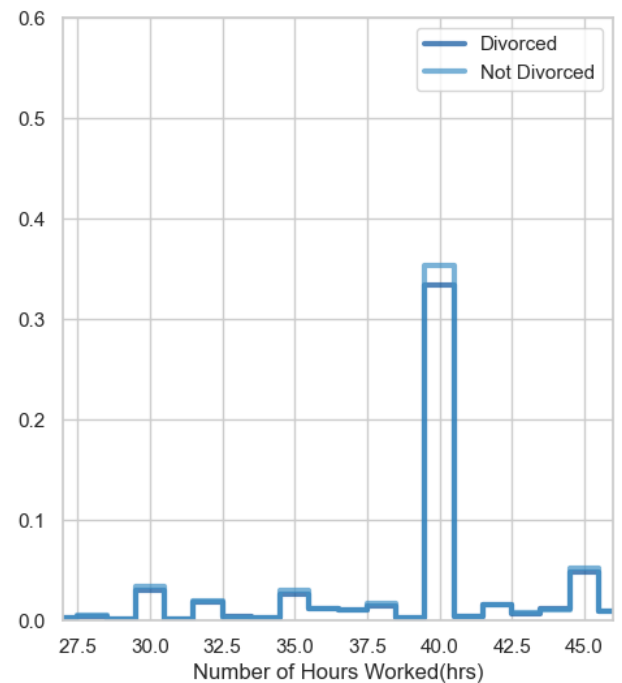

Two scenario comparison by using PMF

```
In [57]: ds_divorced = dataset_gss[dataset_gss.divorce == 'YES']
ds_NOTdivorced = dataset_gss[dataset_gss.divorce != 'YES']

first_pmf = thinkstats2.Pmf(ds_divorced['hrs1'], label="Divorced")
other_pmf = thinkstats2.Pmf(ds_NOTdivorced['hrs1'], label="Not Divorced")

width = 0.45
axis = [27, 46, 0, 0.6]
thinkplot.PrePlot(2, cols=2)
thinkplot.Hist(first_pmf, align="right", width=width)
thinkplot.Hist(other_pmf, align="left", width=width)
thinkplot.Config(xlabel="Number of Hours Worked(hrs)", ylabel="PMF", axis=axis)

thinkplot.PrePlot(2)
thinkplot.SubPlot(2)
thinkplot.Pmfs([first_pmf, other_pmf])
thinkplot.Config(xlabel="Number of Hours Worked(hrs)", axis=axis)
```



Calculating PMF and CDF

```
In [58]: def PlotCDFAndPMF(dataset, xlabel):
    cdf = thinkstats2.Cdf(dataset)

    # Calculate PMF
    pmf = thinkstats2.Pmf(dataset)

    # Create a figure with two subplots
    fig, (ax1) = plt.subplots(1, 1, figsize=(8, 5))

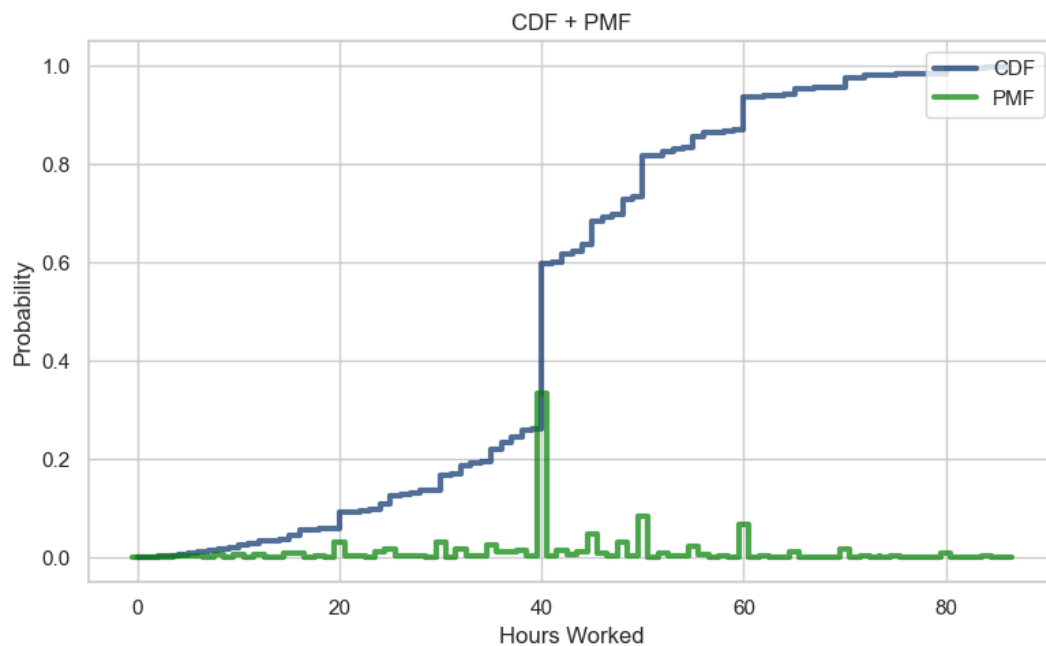
    # Plot CDF on the first subplot
    thinkplot.Cdf(cdf, label='CDF')
    ax1.set_xlabel(xlabel)
    ax1.set_title("CDF + PMF")
    ax1.set_ylabel("Probability")

    # Plot PMF on the second subplot
    thinkplot.Pmf(pmf, label='PMF', color='green')

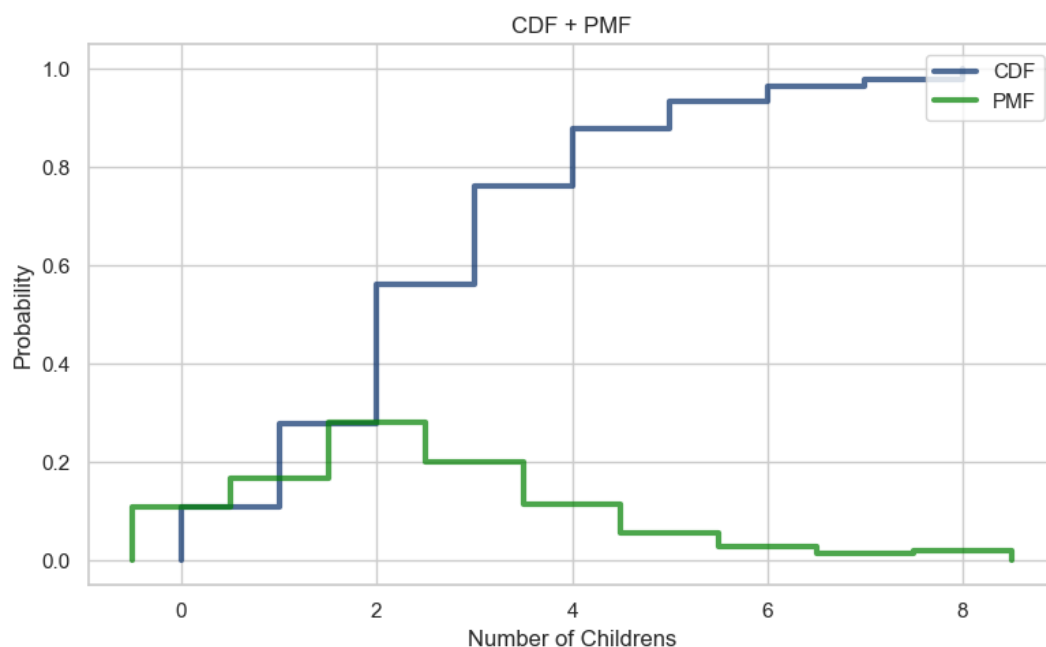
    # Adjust layout and display the plots
    lines, labels = ax1.get_legend_handles_labels()
    #lines2, labels2 = ax2.get_legend_handles_labels()
    ax1.legend(lines, labels, loc="upper right")

    plt.tight_layout()
    plt.show()
```

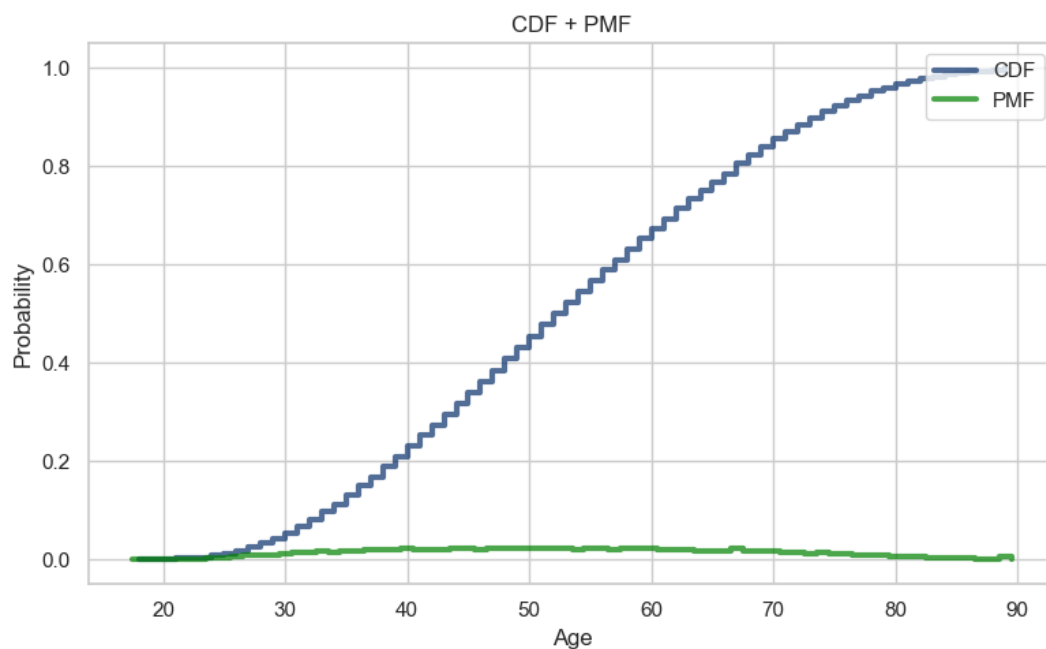
```
In [59]: PlotCDFAndPMF(ds_divorced['hrs1'], "Hours Worked")
```



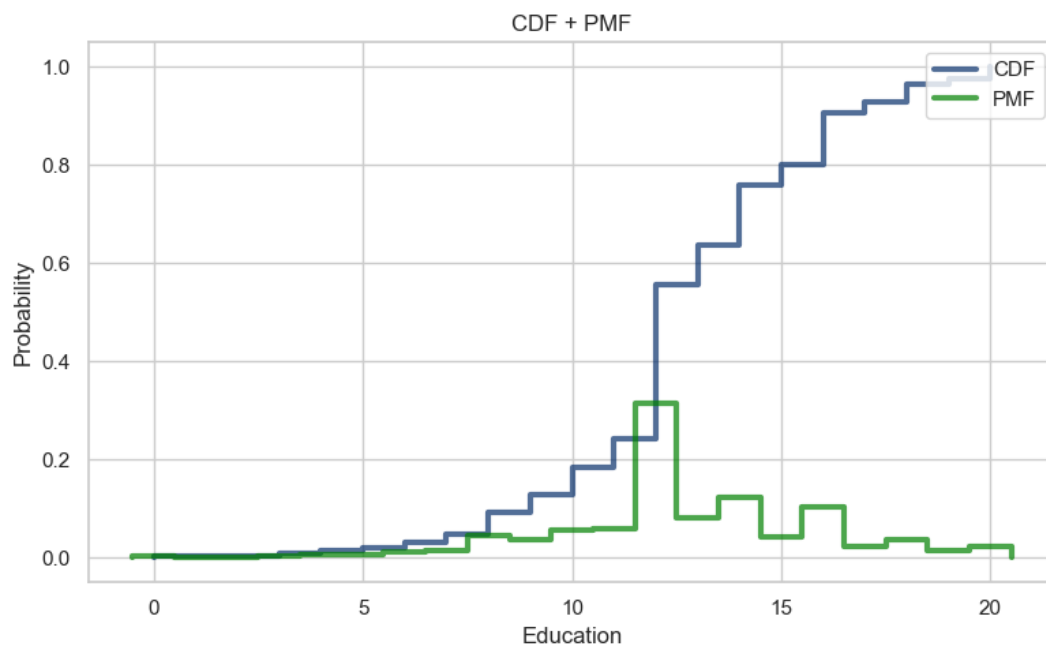
```
In [60]: #['hrs1', 'wrkslf', 'spwrksta', 'chlds', 'age', 'educ', 'sex', 'race', 'family16', 'income']
PlotCDFAndPMF(ds_divorced['chlds'], "Number of Childrens")
#cdf = thinkstats2.Cdf(ds_divorced['chlds'])
#thinkplot.Cdf(cdf)
#thinkplot.Show(xLabel="Number of Childrens")
```



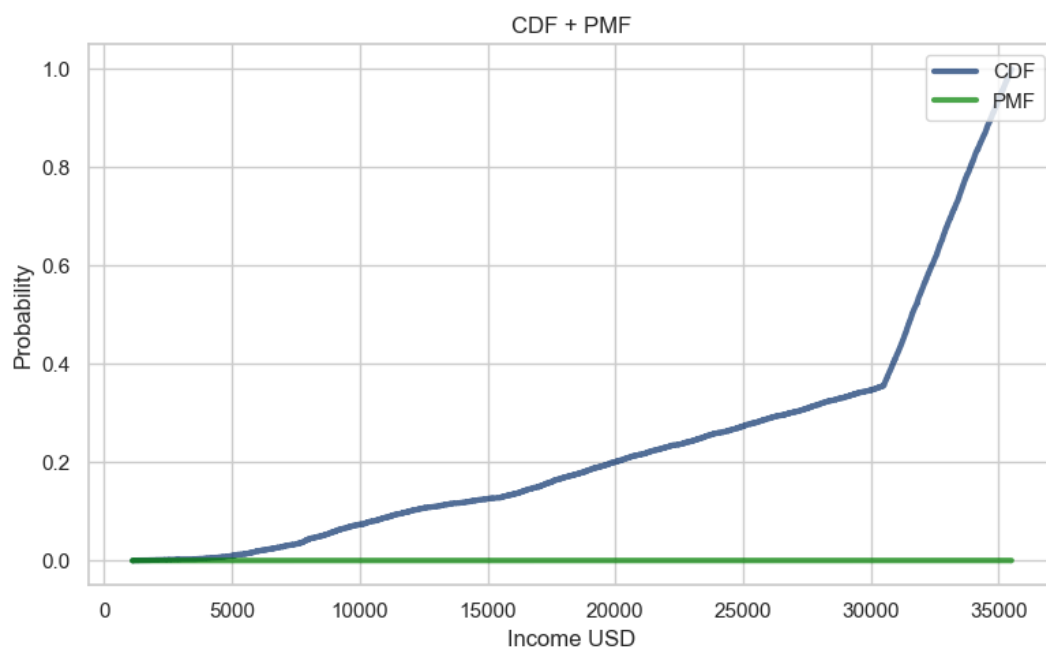
```
In [61]: #['hrs1', 'wrkslf', 'spwrksta', 'chlds', 'age', 'educ', 'sex', 'race', 'family16', 'income']
#cdf = thinkstats2.Cdf(ds_divorced['age'])
#thinkplot.Cdf(cdf)
#thinkplot.Show(xLabel="Age")
PlotCDFAndPMF(ds_divorced['age'], "Age")
```




```
In [62]: #['hrs1', 'wrkslf', 'spwrksta', 'chlds', 'age', 'educ', 'sex', 'race', 'family16', 'income']
PlotCDFAndPMF(ds_divorced['educ'], "Education")
```




```
In [63]: PlotCDFAndPMF(ds_divorced['incomeUSD'], "Income USD")
```



Analytical distribution

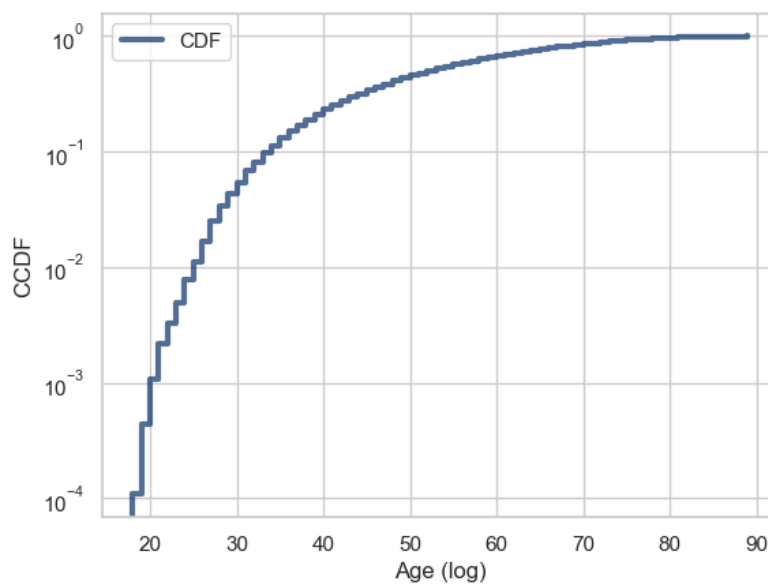
```
In [64]: def PlotCDFAnalytical(dataset, xlabel):
cdf = thinkstats2.Cdf(dataset)
thinkplot.Cdf(cdf, label='CDF')
thinkplot.Show(xlabel=xlabel, ylabel='CCDF', yscale='log')
```

```
In [65]: PlotCDFAnalytical(ds_divorced['hrs1'], "Hours(log)")
```



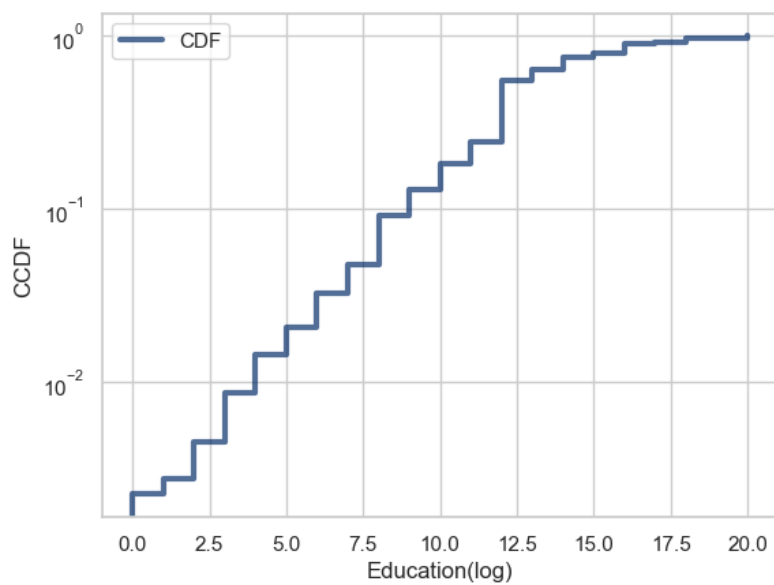
<Figure size 800x600 with 0 Axes>

```
In [66]: PlotCDFAnalytical(ds_divorced['age'], "Age (log)")
```



<Figure size 800x600 with 0 Axes>

```
In [67]: PlotCDFAnalytical(ds_divorced['educ'], "Education(log)")
```



<Figure size 800x600 with 0 Axes>

Scatter plot of different variables

```
In [68]: #Plotting scatter plot of different variables
#the blue shows not divorced, the orange shows the divorce

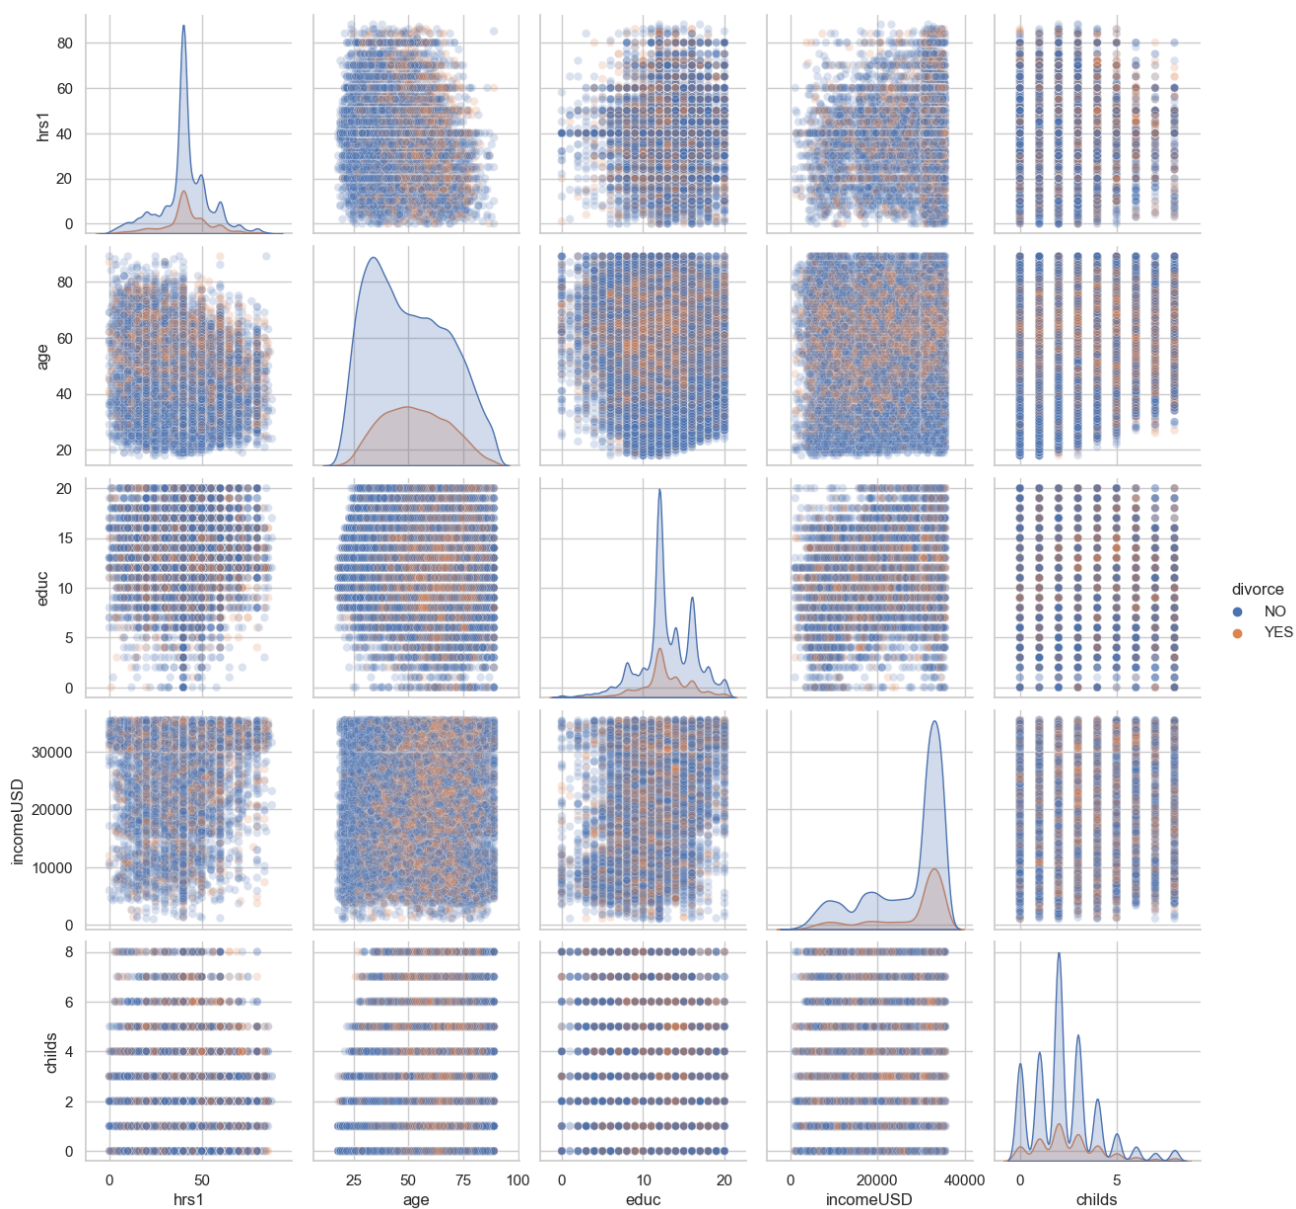
# Load the dataset
dataset = dataset_gss[['divorce', 'hrs1', 'age', 'educ', 'incomeUSD', 'childs', 'spwrksta',
                        'sex', 'race', 'family16', 'wrks1f']]

df = dataset.copy()

# Set the style for the plots (optional, just for aesthetics)
sns.set(style='whitegrid')

# Create a pair plot
sns.pairplot(df, diag_kind='kde', hue='divorce', plot_kws={'alpha': 0.2})

# Show the plot
plt.show()
```



```
In [69]: #Plotting scatter plot of divorce variables

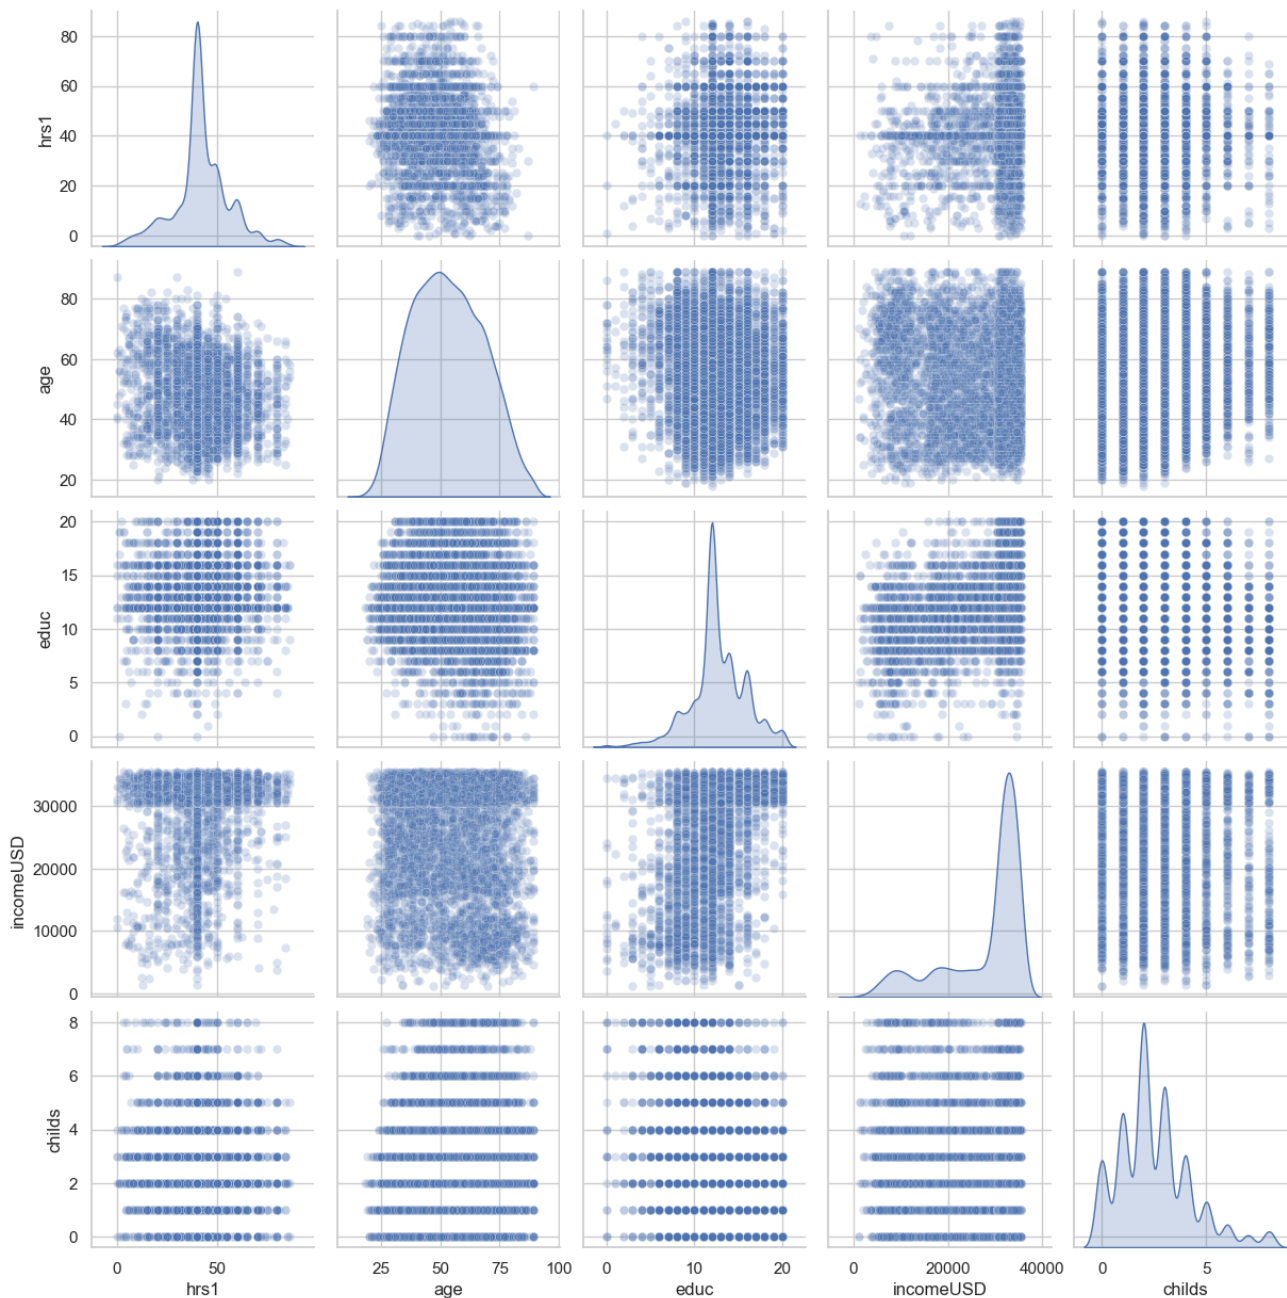
dataset = ds_divorced[['divorce', 'hrs1', 'age', 'educ', 'incomeUSD', 'childs', 'spwrksta',
                      'sex', 'race', 'family16', 'wrks1f']]

df = dataset


# Set the style for the plots (optional, just for aesthetics)
sns.set(style='whitegrid')

# Create a pair plot
sns.pairplot(df, diag_kind='kde', plot_kws={'alpha': 0.2})


# Show the plot
plt.show()
```



```
In [70]: #Hexbin plot, that shows the age and childs relationship for divorced couples
thinkplot.HexBin(ds_divorced['age'],ds_divorced['childs'])
#[['divorce','hrs1','age','educ','incomeUSD','childs','spwrksta',
#   'sex','race','family16','wrkslf']]
```



```
In [71]: #Hexbin plot, that shows the age and childs relationship for divorced couples
thinkplot.HexBin(ds_divorced['age'],ds_divorced['educ'])
#[['divorce','hrs1','age','educ','incomeUSD','childs','spwrksta',
#   'sex','race','family16','wrkslf']]
```

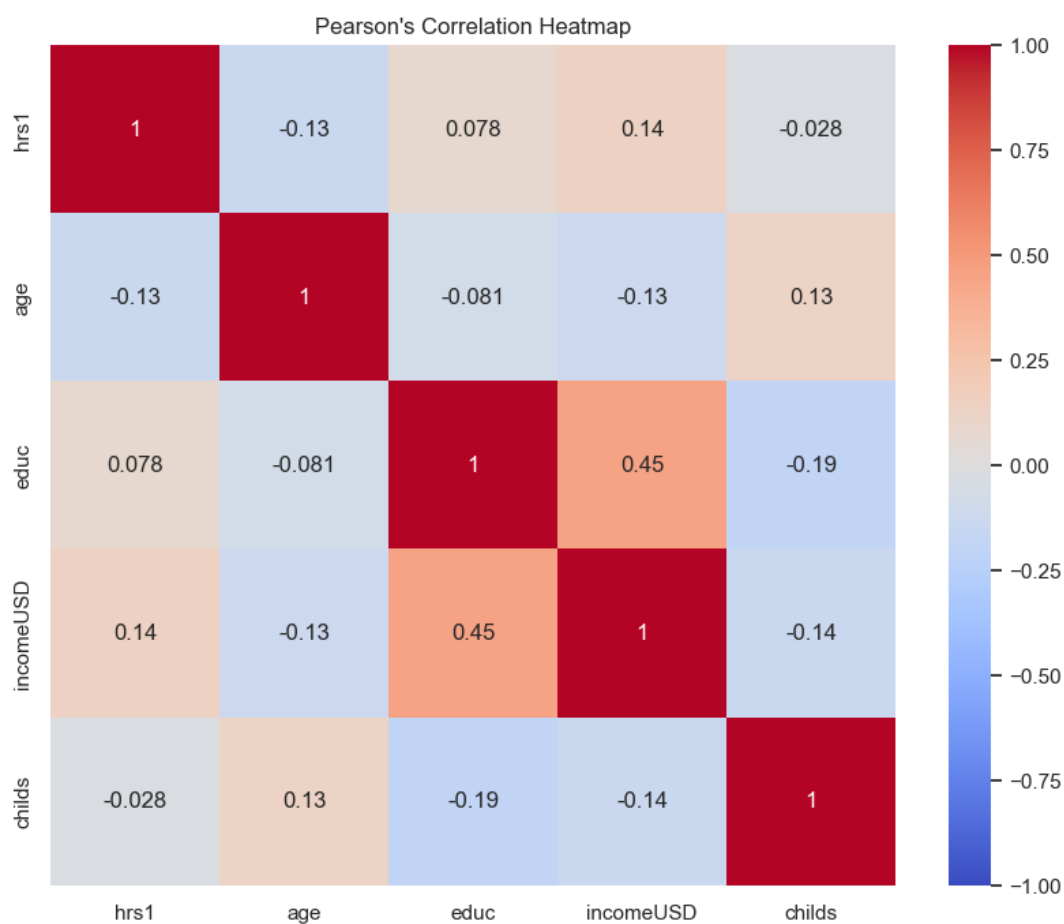


```
In [72]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy.stats import pearsonr
import statsmodels.api as sm

# Load the dataset
dataset = ds_divorced[['hrs1', 'age', 'educ', 'incomeUSD', 'childs']]

# Compute Pearson's correlation matrix
correlation_matrix = dataset.corr()

# Display the correlation matrix as a heatmap
plt.figure(figsize=(10, 8))
sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title("Pearson's Correlation Heatmap")
plt.show()
```



Running OLS regression

```
In [73]: # Perform regression analysis to assess relationships

dataset_analysis=dataset_gss.copy()
# Convert 'divorce' to integer

# Convert 'divorce' to numeric type and fill NaN with a placeholder value
dataset_analysis['divorce'] = pd.to_numeric(dataset_analysis['divorce'], errors='coerce')
dataset_analysis['divorce'] = dataset_analysis['divorce'].fillna(-999).astype(int)

# Perform regression analysis to assess relationships
y = dataset_analysis['divorce']

X = dataset_analysis[['hrs1', 'age', 'educ', 'incomeUSD', 'childs']]
X = X.replace([np.inf, -np.inf], np.nan).fillna(0)
X = sm.add_constant(X) # Add a constant term for the intercept

model = sm.OLS(y, X).fit()
print(model.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          divorce    R-squared:                -inf
Model:                  OLS        Adj. R-squared:             -inf
Method:                 Least Squares    F-statistic:          -1.448e+04
Date:                  Sat, 12 Aug 2023    Prob (F-statistic):      1.00
Time:                  11:50:28          Log-Likelihood:        1.8772e+06
No. Observations:      72390            AIC:                  -3.754e+06
Df Residuals:          72384            BIC:                  -3.754e+06
Df Model:              5
Covariance Type:       nonrobust
=====
                        coef    std err          t      P>|t|      [0.025    0.975]
-----
const          -999.0000    2.65e-14   -3.78e+16    0.000    -999.000    -999.000
hrs1           -2.043e-14    2.42e-16   -84.341     0.000    -2.09e-14    -2e-14
age            1.41e-14     3.01e-16    46.823     0.000     1.35e-14     1.47e-14
educ           1.066e-13     1.65e-15    64.628     0.000     1.03e-13     1.1e-13
incomeUSD      4.033e-17     4.43e-19    90.961     0.000     3.95e-17     4.12e-17
childs         5.063e-14     3.03e-15    16.690     0.000     4.47e-14     5.66e-14
=====
Omnibus:                 35.174    Durbin-Watson:           0.565
Prob(Omnibus):            0.000    Jarque-Bera (JB):        38.698
Skew:                     0.013    Prob(JB):                 3.95e-09
Kurtosis:                 3.110    Cond. No.                 1.39e+05
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 1.39e+05. This might indicate that there are strong multicollinearity or other numerical problems.

```
C:\ProgramData\anaconda3\lib\site-packages\statsmodels\regression\linear_model.py:1752: RuntimeWarning: divide by zero encountered in double_scalars
    return 1 - self.ssr/self.centered_tss
```

The Null Hypothesis

The Null hypothesis 1: There is no co-rellation between divorce and the number of childrens


```
In [74]: # Create a DataFrame
df = dataset_gss.copy()

# Create a contingency table
contingency_table = pd.crosstab(df['divorce'], df['childs'])

# Perform the Chi-Square test
chi2, p, dof, expected = chi2_contingency(contingency_table)

# Define significance level
alpha = 0.05

# Print the results
print("\nChi-Square:", chi2)
print("p-value:", p)
print("Degrees of Freedom:", dof)

# Compare p-value with significance level
if p < alpha:
    print("\nReject the null hypothesis. There is a significant correlation between divorce and the number of children.")
else:
    print("\nFail to reject the null hypothesis. There is no significant correlation between divorce and the number of children.")
```

Chi-Square: 212.3517343739779
 p-value: 1.5873059852030754e-41
 Degrees of Freedom: 8

Reject the null hypothesis. There is a significant correlation between divorce and the number of children.

Reject the null hypothesis. There is a significant correlation between divorce and the number of children.

The Null hypothesis 2: There is no co-rellation between divorce and the number of hours worked in a week

```
In [75]: # Create a DataFrame
df = dataset_gss.copy()

# Create a contingency table
contingency_table = pd.crosstab(df['divorce'], df['hrs1'])

# Perform the Chi-Square test
chi2, p, dof, expected = chi2_contingency(contingency_table)

# Define significance level
alpha = 0.05

# Print the results
print("\nChi-Square:", chi2)
print("p-value:", p)
print("Degrees of Freedom:", dof)

# Compare p-value with significance level
if p < alpha:
    print("\nReject the null hypothesis. There is a significant correlation between divorce and the number of hours worked.")
else:
    print("\nFail to reject the null hypothesis. There is no significant correlation between divorce and the number of hours worked.")
```

Chi-Square: 103.43880050111484
 p-value: 0.12476633823367769
 Degrees of Freedom: 88

Fail to reject the null hypothesis. There is no significant correlation between divorce and the number of hours worked.

Fail to reject the null hypothesis. There is no significant correlation between divorce and the number of hours worked.

Logistic Regression

```
In [76]: # Check if the column is numeric
def is_numeric(column):
    return pd.api.types.is_numeric_dtype(column)

# Impute the missing values
def impute_missing_values(df):
    for column in df.columns:
        if is_numeric(column):
            imputer = SimpleImputer(missing_values='NaN', strategy='mean')
            df[column] = imputer.fit_transform(df[column].values.reshape(-1, 1))
        else:
            #df[column] = df[column].fillna('missing')
            a=2

# Impute the missing values in the dataset
```

```
In [77]: dataset_gss
```

Out[77]:

	year	id_	hrs1	wrkslf	marital	divorce	spwrksta	childs	age	educ	spdeg	sex	race	family16	income	relig16	hapmar	b
0	1972	1	NaN	Someone else	Never married	NaN	NaN	0.0	23.0	16.0	NaN	FEMALE	White	FATHER	NaN	NaN	NaN	<
1	1972	2	NaN	Someone else	Married	NO	KEEPING HOUSE	5.0	70.0	10.0	HIGH SCHOOL	MALE	White	M AND F RELATIVES	NaN	NaN	NaN	<
2	1972	3	NaN	Someone else	Married	NO	WORKING FULLTIME	4.0	48.0	12.0	NaN	FEMALE	White	MOTHER & FATHER	NaN	NaN	NaN	<
3	1972	4	NaN	Someone else	Married	NO	WORKING FULLTIME	0.0	27.0	17.0	GRADUATE	FEMALE	White	MOTHER & FATHER	NaN	NaN	NaN	<
4	1972	5	NaN	Someone else	Married	NO	TEMP NOT WORKING	2.0	61.0	12.0	HIGH SCHOOL	FEMALE	White	MOTHER & FATHER	NaN	NaN	NaN	<
...
72385	2022	3541	48.0	Someone else	Never married	NaN	NaN	0.0	22.0	12.0	NaN	FEMALE	White	MOTHER & STPFATHER	\$25,000 or more	NaN	NaN	B
72386	2022	3542	50.0	Someone else	Married	YES	NaN	2.0	29.0	19.0	NaN	FEMALE	White	MOTHER & FATHER	\$25,000 or more	NaN	VERY HAPPY	B
72387	2022	3543	38.0	Someone else	Never married	NaN	NaN	1.0	32.0	15.0	NaN	MALE	White	MOTHER & STPFATHER	\$25,000 or more	NaN	NaN	B
72388	2022	3544	40.0	Someone else	Married	NO	NaN	0.0	49.0	17.0	NaN	FEMALE	White	MOTHER & FATHER	\$25,000 or more	NaN	VERY HAPPY	B
72389	2022	3545	40.0	Someone else	Married	NO	NaN	1.0	50.0	20.0	NaN	MALE	White	MOTHER & FATHER	\$25,000 or more	NaN	PRETTY HAPPY	B

72390 rows × 19 columns

```
In [78]: # Create the dependent variable
df = dataset_gss.copy()
# Create the independent variables
df = df.drop('id_', axis=1)
df = df.drop('ballot', axis=1)
df = df.drop('year', axis=1)
#df = df.drop('marital', axis=1)

data = df
# Impute NaN values with the mean for numeric columns and the mode for categorical columns
for col in data.columns:
    if data[col].dtype.name == "float64":
        data[col].fillna(data[col].mean(), inplace=True)
    elif not pd.api.types.is_numeric_dtype(data[col]):
        data = pd.get_dummies(data, columns=[col])
        #print(col)

    # Split the data into training and testing sets
data = data.drop('divorce_NO', axis=1)
#X_train, X_test, y_train, y_test = train_test_split(data.drop(columns="marital_Divorced"), data["marital_Divorced"], test_size=0.25)
X_train, X_test, y_train, y_test = train_test_split(data.drop(columns="divorce_YES"), data["divorce_YES"], test_size=0.25)

# Create a Logistic regression model
log_reg = LogisticRegression(max_iter = 10000)

# Fit the model to the training data
result = log_reg.fit(X_train, y_train)

# Make predictions on the testing data
predictions = log_reg.predict(X_test)

# Calculate the accuracy of the model
accuracy = np.mean(predictions == y_test)

print("The accuracy of the model is:", accuracy)
```

The accuracy of the model is: 0.8681622278704829

```

In [79]: import statsmodels.api as sm
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import Ridge
# Scale the input features using StandardScaler
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# Create a Ridge regression model with alpha parameter for regularization
ridge_reg = Ridge(alpha=1.0)
ridge_reg.fit(X_train_scaled, y_train)

# Print coefficients and summary
#print("Ridge Regression Coefficients:")
#for col, coef in zip(X_train.columns, ridge_reg.coef_):
#    print(f"{col}: {coef:.4f}")

# Using statsmodels to fit the Ridge regression model and print summary
X_train_scaled_with_const = sm.add_constant(X_train_scaled)
ridge_reg_sm = sm.OLS(y_train, X_train_scaled_with_const)
result_sm = ridge_reg_sm.fit()

# Print the summary including coefficients and p-values
print(result_sm.summary())

# Make predictions on the testing data
X_test_scaled_with_const = sm.add_constant(X_test_scaled)
y_pred = result_sm.predict(X_test_scaled_with_const)
y_pred_binary = np.round(y_pred) # Convert predicted probabilities to binary predictions

# Calculate accuracy
accuracy = np.mean(y_pred_binary == y_test)
print("Accuracy:", accuracy)

# Calculate precision
precision = np.sum((y_pred_binary == 1) & (y_test == 1)) / np.sum(y_pred_binary == 1)
print("Precision:", precision)

# List column names with their p-values
#p_values = result_sm.pvalues[1:] # Exclude the constant term
#print("\nColumn Names with P-values:")
#for col, p_value in zip(X_train.columns, p_values):
#    print(f"{col}: p-value={p_value:.4f}")

```

OLS Regression Results

```

=====
Dep. Variable:      divorce_YES      R-squared:      0.116
Model:              OLS              Adj. R-squared: 0.115
Method:             Least Squares    F-statistic:    106.3
Date:               Sat, 12 Aug 2023  Prob (F-statistic): 0.00
Time:               11:50:30          Log-Likelihood: -14430.
No. Observations:   54292            AIC:            2.900e+04
Df Residuals:       54224            BIC:            2.960e+04
Df Model:           67
Covariance Type:    nonrobust
=====

```

	coef	std err	t	P> t	[0.025	0.975]
const	0.1295	0.001	95.526	0.000	0.127	0.132
x1	0.0055	0.001	3.956	0.000	0.003	0.008
x2	0.0133	0.002	8.299	0.000	0.010	0.016
x3	0.0211	0.002	11.360	0.000	0.017	0.025
x4	-0.0081	0.002	-4.875	0.000	-0.011	-0.005
x5	0.0076	0.009	0.859	0.391	-0.010	0.025
x6	0.0207	0.002	9.018	0.000	0.016	0.025
x7	0.0195	0.002	8.553	0.000	0.015	0.024
x8	-0.0203	0.017	-1.196	0.232	-0.053	0.013
x9	0.0954	0.026	3.693	0.000	0.045	0.146
x10	-0.0082	0.021	-0.392	0.695	-0.049	0.033
x11	-0.0100	0.009	-1.095	0.273	-0.028	0.008
x12	0.0425	0.015	2.902	0.004	0.014	0.071
x13	-0.0157	0.005	-3.204	0.001	-0.025	-0.006
x14	0.0069	0.002	3.282	0.001	0.003	0.011
x15	-0.0105	0.004	-2.377	0.017	-0.019	-0.002
x16	-0.0020	0.002	-1.095	0.274	-0.006	0.002
x17	-1.436e-05	0.002	-0.007	0.995	-0.004	0.004
x18	0.0043	0.002	1.937	0.053	-5.03e-05	0.009
x19	-0.0036	0.008	-0.467	0.641	-0.019	0.011
x20	-0.0056	0.004	-1.493	0.135	-0.013	0.002
x21	-0.0034	0.003	-1.309	0.190	-0.008	0.002
x22	-0.0185	0.004	-4.784	0.000	-0.026	-0.011
x23	-0.0121	0.003	-3.918	0.000	-0.018	-0.006
x24	-0.0047	0.006	-0.775	0.438	-0.016	0.007
x25	-0.0121	0.004	-2.912	0.004	-0.020	-0.004
x26	-0.0527	0.017	-3.032	0.002	-0.087	-0.019
x27	-0.0508	0.017	-2.916	0.004	-0.085	-0.017
x28	0.0016	0.012	0.135	0.893	-0.022	0.026
x29	-0.0024	0.008	-0.282	0.778	-0.019	0.014
x30	0.0044	0.014	0.314	0.754	-0.023	0.032
x31	-0.0005	0.002	-0.219	0.826	-0.005	0.004
x32	0.0004	0.002	0.209	0.835	-0.004	0.005
x33	0.0007	0.002	0.328	0.743	-0.003	0.005
x34	0.0028	0.002	1.255	0.209	-0.002	0.007
x35	-0.0008	0.002	-0.509	0.611	-0.004	0.002
x36	0.0030	0.004	0.688	0.492	-0.006	0.011
x37	-0.0101	0.006	-1.751	0.080	-0.021	0.001
x38	0.0073	0.003	2.409	0.016	0.001	0.013
x39	0.0100	0.002	4.257	0.000	0.005	0.015
x40	0.0037	0.003	1.171	0.242	-0.002	0.010
x41	0.0007	0.003	0.249	0.803	-0.005	0.007
x42	0.0028	0.002	1.476	0.140	-0.001	0.006
x43	0.0016	0.002	0.907	0.364	-0.002	0.005
x44	0.0156	0.004	3.675	0.000	0.007	0.024
x45	0.0038	0.003	1.275	0.202	-0.002	0.010
x46	0.0007	0.003	0.254	0.799	-0.005	0.006
x47	0.0009	0.003	0.337	0.736	-0.004	0.006
x48	0.0030	0.003	1.194	0.232	-0.002	0.008
x49	0.0009	0.002	0.347	0.728	-0.004	0.006
x50	0.0007	0.003	0.259	0.796	-0.005	0.006
x51	0.0039	0.003	1.382	0.167	-0.002	0.010
x52	-0.0004	0.001	-0.308	0.758	-0.003	0.002
x53	-0.0068	0.003	-2.290	0.022	-0.013	-0.001
x54	0.0008	0.001	0.588	0.556	-0.002	0.004
x55	-0.0024	0.001	-1.699	0.089	-0.005	0.000
x56	0.0006	0.001	0.404	0.686	-0.002	0.003
x57	-0.0026	0.002	-1.632	0.103	-0.006	0.001
x58	0.0005	0.001	0.391	0.696	-0.002	0.003
x59	0.0009	0.001	0.639	0.523	-0.002	0.004
x60	0.0028	0.002	1.480	0.139	-0.001	0.007
x61	0.0006	0.001	0.445	0.656	-0.002	0.003
x62	-0.0017	0.001	-1.153	0.249	-0.005	0.001
x63	-0.0012	0.001	-0.847	0.397	-0.004	0.002
x64	0.0104	0.003	3.215	0.001	0.004	0.017
x65	0.0100	0.002	6.249	0.000	0.007	0.013
x66	0.0044	0.003	1.511	0.131	-0.001	0.010
x67	0.0008	0.003	0.229	0.819	-0.006	0.008

```

=====
Omnibus:      15858.606      Durbin-Watson:      2.006

```

Prob(Omnibus):	0.000	Jarque-Bera (JB):	34941.170
Skew:	1.740	Prob(JB):	0.00
Kurtosis:	4.825	Cond. No.	60.3

=====

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

Accuracy: 0.8709249640844292

Precision: nan

C:\Users\aniruddha.joshi\AppData\Local\Temp\ipykernel_25488\2104285088.py:37: RuntimeWarning: invalid value encountered in long_scalars

```
precision = np.sum((y_pred_binary == 1) & (y_test == 1)) / np.sum(y_pred_binary == 1)
```

```
In [80]: from tabulate import tabulate

# ... (previous code) ...

# List column names with their p-values, coefficients, and R-squared
p_values = result_sm.pvalues[1:] # Exclude the constant term
coefficients = result_sm.params[1:] # Exclude the constant term
r_squared = result_sm.rsquared

# Prepare data for tabular format
table_data = []
for col, coef, p_value in zip(X_train.columns, coefficients, p_values):
    table_data.append([col, coef, p_value])

# Add R-squared to the table
table_data.append(['R-squared', r_squared, ''])

# Print column names, coefficients, p-values, and R-squared in tabular format
headers = ["Column Name", "Coefficient", "P-value"]
print("\nSummary:")
print(tabulate(table_data, headers=headers, floatfmt=(".4f", ".4f", ".4f")))
```

Summary: Column Name	Coefficient	P-value
hrs1	0.0055	7.620459765753077e-05
childs	0.0133	1.0707931742194831e-16
age	0.0211	7.170395867519049e-30
educ	-0.0081	1.0892364007125017e-06
incomeUSD	0.0076	0.39057106719663826
wrkslf_Self-employed	0.0207	1.969026992589378e-19
wrkslf_Someone else	0.0195	1.2257570155047025e-17
marital_Divorced	-0.0203	0.23154110157380986
marital_Married	0.0954	0.00022199688382626233
marital_Never married	-0.0082	0.6954071549873602
marital_Separated	-0.0100	0.2734943796529368
marital_Widowed	0.0425	0.003707304727193523
spwrksta_KEEPING HOUSE	-0.0157	0.001355395561627798
spwrksta_OTHER	0.0069	0.0010331660646010002
spwrksta_RETIRED	-0.0105	0.017449143709530247
spwrksta_SCHOOL	-0.0020	0.2737100273559191
spwrksta_TEMP NOT WORKING	-0.0000	0.9946821541930824
spwrksta_UNEMPL, LAID OFF	0.0043	0.052722356874633225
spwrksta_WORKING FULLTIME	-0.0036	0.6407888968926208
spwrksta_WORKING PARTTIME	-0.0056	0.13543520459486624
spdeg_ASSOCIATE/JUNIOR COLLEGE	-0.0034	0.19043024829369867
spdeg_BACHELOR	-0.0185	1.7202075529369903e-06
spdeg_GRADUATE	-0.0121	8.932172580960702e-05
spdeg_HIGH SCHOOL	-0.0047	0.4382041232576087
spdeg_LT HIGH SCHOOL	-0.0121	0.0035876637112035345
sex_FEMALE	-0.0527	0.002428176600103502
sex_MALE	-0.0508	0.0035514124097361945
race_Black	0.0016	0.8926920953041122
race_Other	-0.0024	0.7777678993970554
race_White	0.0044	0.7537838933933967
family16_FATHER	-0.0005	0.8262792298330345
family16_FATHER & STPMOTHER	0.0004	0.8348212956813824
family16_FEMALE RELATIVE	0.0007	0.742852189237496
family16_M AND F RELATIVES	0.0028	0.20943306531849945
family16_MALE RELATIVE	-0.0008	0.610893507412598
family16_MOTHER	0.0030	0.4916856459009773
family16_MOTHER & FATHER	-0.0101	0.07992793549019443
family16_MOTHER & STPFATHER	0.0073	0.015978661101395963
family16_OTHER	0.0100	2.072576362731709e-05
income_\$1,000 to \$2,999	0.0037	0.24172531458132138
income_\$10,000 to \$14,999	0.0007	0.8032990284171928
income_\$15,000 to \$19,999	0.0028	0.1400385808857833
income_\$20,000 to \$24,999	0.0016	0.36441982611061063
income_\$25,000 or more	0.0156	0.0002382057572925851
income_\$3,000 to \$3,999	0.0038	0.2023103669463342
income_\$4,000 to \$4,999	0.0007	0.7994139869948655
income_\$5,000 to \$5,999	0.0009	0.7358917650276431
income_\$6,000 to \$6,999	0.0030	0.2324851938528776
income_\$7,000 to \$7,999	0.0009	0.7283643349741542
income_\$8,000 to \$9,999	0.0007	0.79563483944189
income_Under \$1,000	0.0039	0.1669678928170715
relig16_BUDDHISM	-0.0004	0.7584393042582245
relig16_CATHOLIC	-0.0068	0.021997904747485846
relig16_CHRISTIAN	0.0008	0.5562945876407357
relig16_HINDUISM	-0.0024	0.08935492394359891
relig16_INTER-NONDENOMINATIONAL	0.0006	0.6859964983459588
relig16_JEWISH	-0.0026	0.10274141372535493
relig16_MUSLIM/ISLAM	0.0005	0.6955463804879863
relig16_NATIVE AMERICAN	0.0009	0.5225048231474876
relig16_NONE	0.0028	0.13895096742026966
relig16_ORTHODOX-CHRISTIAN	0.0006	0.6560301866995284
relig16_OTHER	-0.0017	0.24892654177831486
relig16_OTHER EASTERN	-0.0012	0.3971148018798719
relig16_PROTESTANT	0.0104	0.0013031136007823194
hapmar_NOT TOO HAPPY	0.0100	4.15132086125714e-10
hapmar_PRETTY HAPPY	0.0044	0.13066791693248
hapmar_VERY HAPPY	0.0008	0.818735067547654
R-squared	0.1161	

In []: